

# pyrcn\_tutorial

December 12, 2022

## 1 Building blocks of Reservoir Computing

```
[ ]: from sklearn.datasets import make_blobs

# Generate a toy dataset
U, y = make_blobs(n_samples=100, n_features=10)
```

### 1.1 Input-to-Node

The “Input-to-Node” component describes the connections from the input features to the reservoir and the activation functions of the reservoir neurons. Normally, the input weight matrix  $\mathbf{W}^{\text{in}}$  has the dimension of  $N^{\text{res}} \times N^{\text{in}}$ , where  $N^{\text{res}}$  and  $N^{\text{in}}$  are the size of the reservoir and dimension of the input feature vector  $\mathbf{u}[n]$  with the time index  $n$ , respectively. With

$$\mathbf{r}'[n] = f'(\mathbf{W}^{\text{in}}\mathbf{u}[n] + \mathbf{w}^{\text{bi}}) ,$$

we can describe the non-linear projection of the input features  $\mathbf{u}[n]$  into the high-dimensional reservoir space  $\mathbf{r}'[n]$  via the non-linear input activation function  $f'(\cdot)$ .

The values inside the input weight matrix are usually initialized randomly from a uniform distribution on the interval  $[-1, 1]$  and are afterwards scaled using the input scaling factor  $\alpha_{\text{u}}$ . Since in case of a high dimensional input feature space and/or large reservoir sizes  $N^{\text{res}}$ , this leads to a huge input weight matrix and expensive computations to feed the feature vectors into the reservoir, it was shown that it is sufficient to have only a very small number of connections from the input nodes to the nodes inside the reservoir. Each node of the reservoir may therefore be connected to only  $K^{\text{in}}$  ( $\ll N^{\text{in}}$ ) randomly selected input entries. This makes  $\mathbf{W}^{\text{in}}$  typically very sparse and feeding the feature vectors into the reservoir potentially more efficient.

The bias weights  $\mathbf{w}^{\text{bi}}$  with dimension  $N^{\text{res}}$  are typically initialized by fixed random values from a uniform distribution between  $\pm 1$  and multiplied by the hyper-parameter  $\alpha_{\text{bi}}$ .

```
[ ]: from pyrcn.base.blocks import InputToNode
```

```
[ ]: #
#      / - - - - - /
# ----/ Input-to-Node /-----
# u[n] / _ _ _ _ _ / r'[n]
```

```
# U                                R_i2n

input_to_node = InputToNode(hidden_layer_size=50,
                             k_in=5, input_activation="tanh",
                             input_scaling=1.0, bias_scaling=0.1)
```

```
[ ]: %%time

R_i2n = input_to_node.fit_transform(U)
print(U.shape, R_i2n.shape)
```

```
[ ]: %%timeit

R_i2n = input_to_node.fit_transform(U)
```

## 1.2 Node-to-Node

The “Node-to-Node” component describes the connections inside the reservoir. The output of “Input-to-Node”  $\mathbf{r}'[n]$  together with the output of “Node-to-Node” from the previous time step  $\mathbf{r}[n-1]$  are used to compute the new output of “Node-to-Node”  $\mathbf{r}[n]$  using

$$\mathbf{r}[n] = (1 - \lambda)\mathbf{r}[n-1] + \lambda f(\mathbf{r}'[n] + \mathbf{W}^{\text{res}}\mathbf{r}[n-1]) ,$$

which is a leaky integration of the time-dependent reservoir states  $\mathbf{r}[n]$ .  $f(\cdot)$  acts as the non-linear reservoir activation functions of the neurons in “Node-to-Node”. The leaky integration is equivalent to a first-order lowpass filter. Depending on the leakage  $\lambda \in (0, 1]$ , the reservoir states are globally smoothed.

The reservoir weight matrix  $\mathbf{W}^{\text{res}}$  is a square matrix of the size  $N^{\text{res}}$ . These weights are typically initialized from a standard normal distribution. The Echo State Property (ESP) requires that the states of all reservoir neurons need to decay in a finite time for a finite input pattern. In order to fulfill the ESP, the reservoir weight matrix is typically normalized by its largest absolute eigenvalue and rescaled to a spectral radius  $\rho$ , because it was shown that the ESP holds as long as  $\rho \leq 1$ . The spectral radius and the leakage together shape the temporal memory of the reservoir. Similar as for “Input-to-Node”, the reservoir weight matrix gets huge in case of large reservoir sizes  $N^{\text{res}}$ , it can be sufficient to only connect each node in the reservoir only to  $K^{\text{rec}} (\ll N^{\text{res}})$  randomly selected other nodes in the reservoir, and to set the remaining weights to zero.

To incorporate some information from the future inputs, bidirectional RCNs have been introduced.

```
[ ]: from pyrcn.base.blocks import NodeToNode
```

```
[ ]: #
#      /      /      /      /
# ----/ Input-to-Node /-----/ Node-to-Node /-----
# u[n] / _ _ _ _ _ | r'[n] | _ _ _ _ _ | r[n]
# U      R_i2n      R_n2n
```

```
# Initialize, fit and apply NodeToNode
node_to_node = NodeToNode(hidden_layer_size=50,
                           reservoir_activation="tanh",
                           spectral_radius=1.0, leakage=0.9,
                           bidirectional=False)
```

```
[ ]: %%time

R_n2n = node_to_node.fit_transform(R_i2n)
print(U.shape, R_n2n.shape)
```

```
[ ]: %%timeit

R_n2n = node_to_node.fit_transform(R_i2n)
```

The “Node-to-Output” component is the mapping of the reservoir state  $\mathbf{r}[n]$  to the output  $\mathbf{y}[n]$  of the network. In conventional RCNs, this mapping is trained using (regularized) linear regression. To that end, all reservoir states  $\mathbf{r}[n]$  are concatenated into the reservoir state collection matrix  $\mathbf{R}$ . As linear regression usually contains an intercept term, every reservoir state  $\mathbf{r}[n]$  is expanded by a constant of 1. All desired outputs  $\mathbf{d}[n]$  are collected into the desired output collection matrix  $\mathbf{D}$ . Then, the mapping matrix  $\mathbf{W}^{\text{out}}$  can be computed using

$$\mathbf{W}^{\text{out}} = (\mathbf{R}\mathbf{R}^T + \epsilon\mathbf{I})^{-1} (\mathbf{D}\mathbf{R}^T),$$

where  $\epsilon$  is a regularization parameter.

The size of the output weight matrix  $N^{\text{out}} \times (N^{\text{res}} + 1)$  or  $N^{\text{out}} \times (2 \times N^{\text{res}} + 1)$  in case of a bidirectional “Node-to-Node” determines the total number of free parameters to be trained in the neural network.

After training, the output  $\mathbf{y}[n]$  can be computed using Equation

$$\mathbf{y}[n] = \mathbf{W}^{\text{out}}\mathbf{r}[n].$$

Note that, in general, other training methodologies could be used to compute output weights.

```
[ ]: from sklearn.linear_model import Ridge
```

```
[ ]: #
#      / - - - - - /      / - - - - - /      / - - - - - /
# ----/Input-to-Node /-----/Node-to-Node /-----/Node-to-Output /
# u[n] / - - - - - /r'[n] / - - - - - /r[n] / - - - - - /
# U      R_i2n      R_n2n      /
#                               /
#                               /
#                               y[n] / y_pred

# Initialize, fit and apply NodeToOutput
y_pred = Ridge().fit(R_n2n, y).predict(R_n2n)
```

```
print(y_pred.shape)
```

### 1.3 Predict the Mackey-Glass equation

Set up and train vanilla RCNs for predicting the Mackey-Glass time series with the same settings as used to introduce ESNs. The minimum working example shows the simplicity of implementing a model with PyRCN and the inter-operability with scikit-learn; it needs only four lines of code to load the Mackey-Glass dataset that is part of PyRCN and only two lines to fit the different RCN models, respectively. Instead of the default incremental regression, we have customized the `ELMRegressor()` by using `Ridge` from scikit-learn.

```
[ ]: from sklearn.linear_model import Ridge as skRidge
      from pyrcn.echo_state_network import ESNRegressor
      from pyrcn.extreme_learning_machine import ELMRegressor
      from pyrcn.datasets import mackey_glass

[ ]: # Load the dataset
X, y = mackey_glass(n_timesteps=5000)
# Define Train/Test lengths
trainLen = 1900
X_train, y_train = X[:trainLen], y[:trainLen]
X_test, y_test = X[trainLen:], y[trainLen:]

# Initialize and train an ELMRegressor and an ESNRegressor
esn = ESNRegressor().fit(X=X_train.reshape(-1, 1), y=y_train)
elm = ELMRegressor(regressor=skRidge()).fit(X=X_train.reshape(-1, 1), y=y_train)

print("Fitted models")

[ ]: import matplotlib.pyplot as plt
      import seaborn as sns

      %matplotlib inline

[ ]: fig, axs = plt.subplots()

      sns.lineplot(x=list(range(len(y_test))), y=y_test, ax=axs)
      sns.lineplot(x=list(range(len(y_test))), y=esn.predict(X_test.reshape(-1, 1)),
                    ↪ax=axs)
      sns.lineplot(x=list(range(len(y_test))), y=elm.predict(X_test.reshape(-1, 1)),
                    ↪ax=axs)
```

## 2 Build Reservoir Computing Networks with PyRCN

By combining the building blocks introduced above, a vast number of different RCNs can be constructed. In this section, we build two important variants of RCNs, namely ELMs and ESNs.

### 3 Extreme Learning Machines

The vanilla ELM as a single-layer feedforward network consists of an “Input-to-Node” and a “Node-to-Output” module and is trained in two steps:

1. Compute the high-dimensional reservoir states  $\mathbf{R}'$ , which is the collection of reservoir states  $\mathbf{r}'[n]$ .
2. Compute the output weights  $\mathbf{W}^{\text{out}}$  with  $\mathbf{R}'$ .

```
[ ]: U, y = make_blobs(n_samples=100, n_features=10)
      from pyrcn.extreme_learning_machine import ELMRegressor
```

```
[ ]: # Vanilla ELM for regression tasks with input_scaling
#
#      | - - - - - |      | - - - - - |
# ----|Input-to-Node|----|Node-to-Output|-----
# u[n]| - - - - - |r'[n]| - - - - - |y[n]
#                                     y_pred
#
vanilla_elm = ELMRegressor(input_scaling=0.9)
vanilla_elm.fit(U, y)
print(vanilla_elm.predict(U))
```

Example of how to construct an ELM with a BIP “Input-to-Node” ELMs with PyRCN.

```
[ ]: from pyrcn.base.blocks import BatchIntrinsicPlasticity
```

```
[ ]: # Custom ELM with BatchIntrinsicPlasticity
#
#      | - - - - - |      | - - - - - |
# ----|      BIP      |----|Node-to-Output|-----
# u[n]| - - - - - |r'[n]| - - - - - |y[n]
#                                     y_pred
#
bip_elm = ELMRegressor(input_to_node=BatchIntrinsicPlasticity(),
                       regressor=Ridge(alpha=1e-5))

bip_elm.fit(U, y)
print(bip_elm.predict(U))
```

Hierarchical or Ensemble ELMs can then be built using multiple “Input-to-Node” modules in parallel or in a cascade. This is possible when using using scikit-learn’s `sklearn.pipeline.Pipeline` (cascading) or `sklearn.pipeline.FeatureUnion` (ensemble).

```
[ ]: from sklearn.pipeline import Pipeline, FeatureUnion
```

```
[ ]: # ELM with cascaded InputToNode and default regressor
#
#      | - - - - - |      | - - - - - |      | - - - - - |
#      |      (bip)      |      |      (base)      |      |
```

```
# -----|Input-to-Node1|-----|Input-to-Node2|-----|Node-to-Output |
# u[n] / _ _ _ _ _ / _ _ _ _ _ / r'[n] / _ _ _ _ _ /
#
#
#
# y[n] | y_pred

i2n = Pipeline([('bip', BatchIntrinsicPlasticity()),
                ('base', InputToNode(bias_scaling=0.1))])
casc_elm = ELMRegressor(input_to_node=i2n).fit(U, y)

# Ensemble of InputToNode with activations
#
#      _ _ _ _ _
#      |         (i)        |
#      |----|Input-to-Node1|----|
#      |   | _ _ _ _ _ |   |
#      |   |         |   |
#      |   |         |   |
# ----o                               r'[n]|Node-to-Output |-----
# u[n] /                               |-----| y[n]
#      |   |         (th)       |   |
#      |---|Input-to-Node2|-----|
#      | _ _ _ _ _ |
#

i2n = FeatureUnion([('i', InputToNode(input_activation="identity")),
                    ('th', InputToNode(input_activation="tanh"))])
ens_elm = ELMRegressor(input_to_node=i2n)
ens_elm.fit(U, y)
print(casc_elm, ens_elm)
```

### 3.1 Echo State Networks

ESNs, as variants of RNNs, consist of an “Input-to-Node”, a “Node-to-Node” and a “Node-to-Output” block and are trained in three steps:

1. Compute the neuron input states  $\mathbf{R}'$ , which is the collection of reservoir states  $\mathbf{r}'[n]$ . Note that here the activation function  $f'(\cdot)$  is typically linear.
2. Compute the reservoir states  $\mathbf{R}$ , which is the collection of reservoir states  $\mathbf{r}[n]$ . Note that here the activation function  $f(\cdot)$  is typically non-linear.
3. Compute the output weights  $\mathbf{W}^{\text{out}}$  using
  1. Linear regression with  $\mathbf{R}$  when considering an ESN.
  2. Backpropagation or other optimization algorithm when considering a CRN or when using an ESN with non-linear outputs.

What follows is an example of how to construct such a vanilla ESN with PyRCN, where the `ESNRegressor` internally passes the input features through “Input-to-Node” and “Node-to-Node”, and fits “Node-to-Output” using `pyrcn.linear_model.IncrementalRegression`.

```
[ ]: from pyrcn.echo_state_network import ESNRegressor
```

```
[ ]: # Vanilla ESN for regression tasks with spectral_radius and leakage
#
#      /-----/-----/-----/-----/
# ----/Input-to-Node /----/Node-to-Node /----/Node-to-Output /
# u[n] / ----- /r'[n] / ----- /r[n] / ----- /
#
#                                     /
#                                     /
#                                     y[n] / y_pred
#
vanilla_esn = ESNRegressor(spectral_radius=1, leakage=0.9)
vanilla_esn.fit(U, y)
print(vanilla_esn.predict(U))
```

As for ELMs, various unsupervised learning techniques can be used to pre-train “Input-to-Node” and “Node-to-Node”.

```
[ ]: from pyrcn.base.blocks import HebbianNodeToNode
```

```
[ ]: # Custom ESN with BatchIntrinsicPlasticity and HebbianNodeToNode
#
#      |          (bip)         |          |          (hebb)         |          |          - - - - - |
# ----|Input-to-Node |-----|Node-to-Node |-----|Node-to-Output |
# u[n]| _ _ _ _ _ _ _ |r'[n]|_ _ _ _ _ _ _ |r[n] | _ _ _ _ _ _ _ |
#
#                                     |
#                                     |
#                                     y[n] | y_pred
#
bip_esn = ESNRegressor(input_to_node=BatchIntrinsicPlasticity(),
                       node_to_node=HebbianNodeToNode(),
                       regressor=Ridge(alpha=1e-5))

bip_esn.fit(U, y)
print(bip_esn.predict(U))
```

The term “Deep ESN” can refer to different approaches of hierarchical ESN architectures:

Example of how to construct a rather complex ESN consisting of two layers. It is built out of two small parallel reservoirs in the first layer and a large reservoir in the second layer.

```
[ ]: # Multilayer ESN
#
#                                     u[n]
#                                     |
#                                     |
#                                     |
#                                     o -----
#                                     /
#      /-----/-----/-----/-----/
#      |         |         |         |         |
#      |         (i)       |         (i)       |
#      |Input-to-Node1|   |Input-to-Node2|
```

```

# | - - - - - / | - - - - - /
# |           |r1'[n]           | r2'[n]
# | - - - - - / | - - - - - /
# | (th)       | | (th)       |
# | Node-to-Node1 | | Node-to-Node2 |
# | - - - - - / | - - - - - /
# |           |r1[n]           | r2[n]
# | - - - - - / | - - - - - /
# |           |           |
# |           |           |
# | - - - - - / | - - - - - /
# |           | Node-to-Node3 |
# | - - - - - / | - - - - - /
# |           |
# |           r3[n] |
# | - - - - - / | - - - - - /
# |           | Node-to-Output |
# | - - - - - / | - - - - - /
# |           |
# |           y[n] |

l1 = Pipeline([('i2n1', InputToNode(hidden_layer_size=100)),
               ('n2n1', NodeToNode(hidden_layer_size=100))])

l2 = Pipeline([('i2n2', InputToNode(hidden_layer_size=400)),
               ('n2n2', NodeToNode(hidden_layer_size=400))])

i2n = FeatureUnion([('l1', l1),
                    ('l2', l2)])
n2n = NodeToNode(hidden_layer_size=500)
layered_esn = ESNRegressor(input_to_node=i2n,
                           node_to_node=n2n)

layered_esn.fit(U, y)
print(layered_esn.predict(U))

```

```

[ ]: import numpy as np
from sklearn.decomposition import PCA

# Multiple small reservoirs with different leakages in parallel
res1 = FeatureUnion([
    ("lambda_0.1",
     Pipeline([('i2n', InputToNode(hidden_layer_size=10)),
               ('n2n', NodeToNode(hidden_layer_size=10,
                                   leakage=0.1))])),

```



```

("lambda_0.2",
 Pipeline([('i2n', InputToNode(hidden_layer_size=10)),
           ('n2n', NodeToNode(hidden_layer_size=10,
                               leakage=0.2))])),
("lambda_0.3",
 Pipeline([('i2n', InputToNode(hidden_layer_size=10)),
           ('n2n', NodeToNode(hidden_layer_size=10,
                               leakage=0.3))])),
("lambda_0.4",
 Pipeline([('i2n', InputToNode(hidden_layer_size=10)),
           ('n2n', NodeToNode(hidden_layer_size=10,
                               leakage=0.4))])),])

pca = PCA(n_components=10)

res2 = Pipeline([("i2n", InputToNode(hidden_layer_size=100)),
                 ("n2n", NodeToNode(hidden_layer_size=100))])

i2n = FeatureUnion([("path1",
                    Pipeline([("res1", res1), ("pca", pca),
                                ("res2", res2)])),
                  ("path2", res1)])

n2n = NodeToNode(spectral_radius=0., leakage=1., hidden_layer_size=100+40,
                 predefined_recurrent_weights=np.eye(40+100))

deep_esn = ESNRegressor(input_to_node=i2n, node_to_node=n2n)
deep_esn.fit(U, y)
print(deep_esn.predict(U))

```

### 3.2 Complex example: Optimize the hyper-parameters of RCNs

Example for a sequential parameter optimization with PyRCN. Therefore, a model with initial parameters and various search steps are defined. Internally, `SequentialSearchCV` will perform the list of optimization steps sequentially.

```

[ ]: from sklearn.metrics import make_scorer
     from sklearn.metrics import mean_squared_error
     from sklearn.model_selection import TimeSeriesSplit
     from sklearn.model_selection import (RandomizedSearchCV,
                                         GridSearchCV)

     from scipy.stats import uniform
     from pyrcn.model_selection import SequentialSearchCV
     from pyrcn.datasets import mackey_glass

```

```

[ ]: # Load the dataset
     X, y = mackey_glass(n_timesteps=5000)

```

```

X_train, X_test = X[:1900], X[1900:]
y_train, y_test = y[:1900], y[1900:]

# Define initial ESN model
esn = ESNRegressor(bias_scaling=0, spectral_radius=0, leakage=1)

# Define optimization workflow
scorer = make_scorer(mean_squared_error, greater_is_better=False)
step_1_params = {'input_scaling': uniform(loc=1e-2, scale=1),
                 'spectral_radius': uniform(loc=0, scale=2)}
kwargs_1 = {'n_iter': 200, 'n_jobs': -1, 'scoring': scorer,
            'cv': TimeSeriesSplit()}
step_2_params = {'leakage': [0.2, 0.4, 0.7, 0.9, 1.0]}
kwargs_2 = {'verbose': 5, 'scoring': scorer, 'n_jobs': -1,
            'cv': TimeSeriesSplit()}

searches = [('step1', RandomizedSearchCV, step_1_params, kwargs_1),
            ('step2', GridSearchCV, step_2_params, kwargs_2)]

# Perform the search
esn_opti = SequentialSearchCV(esn, searches).fit(X_train.reshape(-1, 1),
        ↪ y_train)
print(esn_opti)

```

### 3.3 Programming pattern for sequence processing

This complex use-case requires a serious hyper-parameter tuning. To keep the code example simple, we did not include the optimization in this paper and refer the interested readers to the Jupyter Notebook <sup>1</sup> that was developed to produce these results.

```

[ ]: from sklearn.base import clone
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.model_selection import ParameterGrid
from sklearn.metrics import make_scorer

from pyrcn.echo_state_network import ESNClassifier
from pyrcn.metrics import accuracy_score
from pyrcn.datasets import load_digits

[ ]: # Load the dataset
X, y = load_digits(return_X_y=True, as_sequence=True)
print("Number of digits: {0}".format(len(X)))
print("Shape of digits {0}".format(X[0].shape))
# Divide the dataset into training and test subsets
X_tr, X_te, y_tr, y_te = train_test_split(X, y, test_size=0.2,

```

<sup>1</sup><https://github.com/TUD-STKS/PyRCN/blob/main/examples/digits.ipynb>

```

                                random_state=42)
print("Number of digits in training set: {0}".format(len(X_tr)))
print("Shape of the first digit: {0}".format(X_tr[0].shape))
print("Number of digits in test set: {0}".format(len(X_te)))
print("Shape of the first digit: {0}".format(X_te[0].shape))

# These parameters were optimized using SequentialSearchCV
esn_params = {'input_scaling': 0.05077514155476392,
              'spectral_radius': 1.1817858863764836,
              'input_activation': 'identity',
              'k_in': 5,
              'bias_scaling': 1.6045393364745582,
              'reservoir_activation': 'tanh',
              'leakage': 0.03470266988650412,
              'k_rec': 10,
              'alpha': 3.0786517836196185e-05,
              'decision_strategy': "winner_takes_all"}

b_esn = ESNClassifier(**esn_params)

param_grid = {'hidden_layer_size': [50, 100, 200, 400, 500],
              'bidirectional': [False, True]}

for params in ParameterGrid(param_grid):
    esn_cv = cross_validate(clone(b_esn).set_params(**params),
                            X=X_tr, y=y_tr,
                            scoring=make_scorer(accuracy_score))
    esn = clone(b_esn).set_params(**params).fit(X_tr, y_tr, n_jobs=1)
    acc_score = accuracy_score(y_te, esn.predict(X_te))
    print(acc_score)

```

[ ]: